

# Feature Selection

Bruno Iochins Grisci

Generative AI Academy

Part 2: Benchmark datasets

# Overview

- 1 Tabular Data
- 2 Synthetic Data
- 3 Real-World Data


# Tabular Data

# Why Tabular Data?

- Common format in many real-world domains: healthcare, finance, bioinformatics.
- Each row is a sample; each column is a feature.
- Fixed-length, interpretable, and structured.
- Suitable for statistical and machine learning models.

**Feature Selection is most often applied to tabular data.**

# Tabular data



Sepal length	Sepal width	Petal length	Petal width	Species
5.1	3.5	1.4	0.2	<i>I. setosa</i>
4.9	3.0	1.4	0.2	<i>I. setosa</i>

# Tabular data

- 1 Receives less attention than image and text data from the interpretability community [Molnar, 2019].
- 2 One of the domains not dominated by deep learning.
- 3 But this is changing [Borisov et al., 2021] [Kadra et al., 2021].
- 4 Much of the available data in science and business is tabular [Borisov et al., 2021].
- 5 Surrogate models (LIME) [Ribeiro et al., 2016] require the definition of a neighborhood, which is not well defined for tabular data [Molnar, 2019].
- 6 The attention layer is restricted to categorical features and is not applied to continuous features [Huang et al., 2020].

# Synthetic Data

# XOR Problem

## Motivating example for feature interactions.

- Binary classification problem with non-linear separability.
- Cannot be solved with univariate feature selection.
- Requires multivariate or interaction-aware methods.

$$\text{XOR}(x_1, x_2) = x_1 \oplus x_2$$



# Generating XOR with Random Features

## Steps to create an XOR dataset:

- Sample two binary features  $x_1, x_2 \sim \text{Bernoulli}(0.5)$
- Label is  $y = x_1 \oplus x_2$  (XOR logic)
- Add random noise features (irrelevant)

## XOR

Class	Informative 1	Informative 2	Noisy 1	Noisy 2	Noisy 3	...	Noisy 48
0	0	0	0	1	0	...	1
1	0	1	0	1	0	...	1
1	1	0	1	0	1	...	1
1	1	0	0	0	0	...	1
0	0	0	1	0	0	...	0
1	1	0	0	0	1	...	0
0	0	0	0	1	0	...	0
0	1	1	0	1	1	...	0

# Synthetic Data

## Why use synthetic data for feature selection?

- Controlled experiments with known ground truth.
- Evaluate selection accuracy (PIFS, PSFI).
- Test scalability and robustness.

## Common designs:

- Informative + redundant + irrelevant features
- Binary or continuous variables


# Creating Synthetic Datasets in Python

## Strategy for reproducible and interpretable datasets:

- Define relevant features (informative to the label)
- Generate redundant features (e.g., linear combinations)
- Add irrelevant noise features

```
sklearn.datasets.make_classification(n_samples=100,  
n_features=20, *, n_informative=2, n_redundant=2,  
n_repeated=0, n_classes=2, n_clusters_per_class=2,  
weights=None, flip_y=0.01, class_sep=1.0, hypercube=True,  
shift=0.0, scale=1.0, shuffle=True, random_state=None)
```

# Synthetic Datasets




# Real-World Data

# UCI Machine Learning Repository

## A classic benchmark collection for ML and FS.

- Curated dataset repository since 1987.
- Contains tabular datasets for classification, regression, clustering.
- Well-documented and widely used in research.
- Examples: Iris, Wine, Breast Cancer, Spam, Heart Disease.
- Drawbacks: age and size.
- <https://archive.ics.uci.edu/>

# UCI Machine Learning Repository



## Bank Marketing

Donated on 2/13/2012

The data is related with direct marketing campaigns (phone calls) of a Portuguese banking institution. The classification goal is to predict if the client will subscribe a term deposit (variable y).

Dataset Characteristics	Subject Area	Associated Tasks
Multivariate	Business	Classification

Feature Type	# Instances	# Features
Categorical, Integer	45211	16

[Download](#) (999.0 KB)

[Import in Python](#)

[Cite](#)

9 citations  
431824 views

**Creators**

- S. Moro
- P. Rita
- P. Cortez

**DOI**

10.24432/CSK306

**License**

This dataset is licensed under a [Creative Commons Attribution 4.0 International](#) (CC BY 4.0) license.

This allows for the sharing and adaptation of the datasets for any purpose, provided that the appropriate credit is given.

**Dataset Information**

**Additional Information**

The data is related with direct marketing campaigns of a Portuguese banking institution. The marketing campaigns were based on phone calls. Often, more than one contact to the same client was required, in order to access if the product (bank term deposit) would be ('yes') or not ('no') subscribed. ...

[SHOW MORE](#)

**Has Missing Values?**

No

**Introductory Paper**

[A data-driven approach to predict the success of bank telemarketing](#)

by Sérgio Moro, P. Cortez, P. Rita. 2014  
Published in Decision Support Systems

**Variables Table**

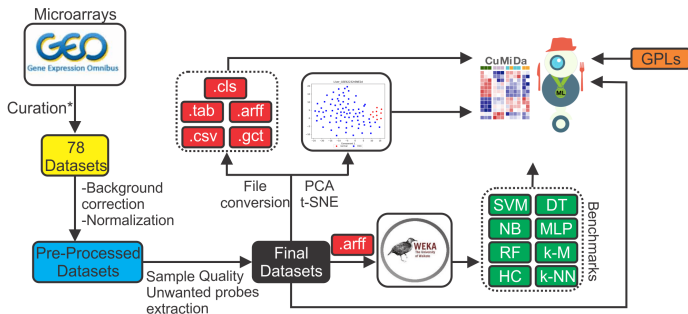


# CuMiDa (Curated Microarray Database)

## **Focus on gene expression microarray data for FS research.**

- Designed for reproducibility and reliability.
- Addresses issues in common microarray datasets.
- Contains metadata: sample origin, platform, preprocessing.
- Facilitates fair comparison of FS algorithms.
- <https://sbc.b.inf.ufgrs.br/cumida>

## CuMiDa



# Being Cautious with Datasets

## Why careful selection and documentation matter.

- Dataset bias can lead to misleading FS conclusions.
- Many popular datasets contain mislabeled or unbalanced data.
- Some are outdated or no longer representative.
- Always document preprocessing, sources, and limitations.

## Critical for reliable FS research and reproducibility.

# References I



Borisov, V., Leemann, T., Seßler, K., Haug, J., Pawelczyk, M., and Kasneci, G. (2021).

Deep neural networks and tabular data: A survey.

*arXiv preprint arXiv:2110.01889.*



Huang, X., Khetan, A., Cvitkovic, M., and Karnin, Z. (2020).

Tabtransformer: Tabular data modeling using contextual embeddings.

*arXiv preprint arXiv:2012.06678.*



Kadra, A., Lindauer, M., Hutter, F., and Grabocka, J. (2021).

Regularization is all you need: Simple neural nets can excel on tabular data.

*arXiv preprint arXiv:2106.11189.*

# References II



Molnar, C. (2019).

*Interpretable Machine Learning.*

<https://christophm.github.io/interpretable-ml-book/>.



Ribeiro, M. T., Singh, S., and Guestrin, C. (2016).

"why should i trust you?" explaining the predictions of any classifier.

In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 1135–1144, San Francisco California USA. ACM.